



EXECUTIVE OFFICE OF THE PRESIDENT
OFFICE OF MANAGEMENT AND BUDGET
WASHINGTON, D.C. 20503

THE DIRECTOR

December 11, 2025

M-26-04

MEMORANDUM TO HEADS OF EXECUTIVE DEPARTMENTS AND AGENCIES

FROM: Russell T. Vought
Director

SUBJECT: Increasing Public Trust in Artificial Intelligence Through Unbiased AI Principles

Overview

On July 23, 2025, President Trump signed Executive Order 14319, *Preventing Woke AI in the Federal Government*¹ (the E.O.), to ensure that Large Language Models (LLMs) procured by the Federal Government produce reliable outputs free from harmful ideological biases or social agendas. To further this policy of promoting innovation and use of trustworthy artificial intelligence (AI), the E.O. identifies two Unbiased AI Principles:²

(1) Truth-seeking

LLMs³ shall be truthful in responding to user prompts seeking factual information or analysis. LLMs shall prioritize historical accuracy, scientific inquiry, and objectivity, and shall acknowledge uncertainty where reliable information is incomplete or contradictory.

(2) Ideological Neutrality

LLMs shall be neutral, nonpartisan tools that do not manipulate responses in favor of ideological dogmas. Developers shall not intentionally encode partisan or ideological judgments into an LLM's outputs unless those judgments are prompted by or otherwise readily accessible to the end user.

Section 4 of the E.O. requires the Director of the Office of Management and Budget (OMB) to issue guidance to agencies to implement these principles. This memorandum fulfills

¹ Executive Order 14139, *Preventing Woke AI in the Federal Government* (July 23, 2025), <https://www.whitehouse.gov/presidential-actions/2025/07/preventing-woke-ai-in-the-federal-government/>

² Id. sec 3

³ Section 2 of the E.O. defines "LLM" to mean, "a large language model, which is a generative AI model trained on vast, diverse datasets that enable the model to generate natural-language responses to user prompts."

that requirement and complements OMB Memorandum M-25-22, *Driving Efficient Acquisition of Artificial Intelligence in Government*.⁴

Scope

The requirements of this memorandum apply to each agency, where “agency” means an executive department, military department, or any independent establishment within the meaning of 5 U.S.C. 101, 102, and 104(1), respectively, or any wholly owned Government corporation within the meaning of 31 U.S.C. 9101. They do not apply to national security systems, as defined in 44 U.S.C. 3552(b)(6).⁵ However, application of this memorandum to national security systems, to the extent practicable, is encouraged.

Within this scope, the requirements of this memorandum apply to any LLM procured by an agency, regardless of the manner in which the LLM will be deployed, further modified, or used by the agency.⁶ Agencies must also consider the relevant factors identified in Appendix A.2 in determining whether to apply these requirements to LLMs developed by the agency, and to AI models other than LLMs that are procured by the agency.⁷

This memorandum does not govern:

- a. Agencies’ regulatory actions designed to prescribe law or policy regarding non-agency uses of AI;
- b. Agencies’ assessments of particular AI applications because the AI provider is the target or potential target of a regulatory enforcement, law enforcement, or national security action; or the agency is evaluating the AI application because it was used by a criminal suspect;
- c. Agencies’ development of metrics, methods, and standards to test and measure AI, where such metrics, methods, and standards are for use by the general public or the Government

⁴ OMB Memorandum M-25-22, *Driving Efficient Acquisition of Artificial Intelligence in Government* (Apr. 3, 2025), <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf>

⁵ AI innovation and risk for national security systems must be managed appropriately, but these systems are governed through other policy. Agencies should reference existing guidelines in place, such as the Department of War’s (DoW) Responsible Artificial Intelligence Strategy and Implementation Pathway and the Office of the Director of National Intelligence’s Principles of Artificial Intelligence Ethics for the Intelligence Community, as well as policies governing specific high-risk national security applications of AI, such as DoD Directive 3000.09, Autonomy in Weapon Systems, <https://www.esd.whs.mil/portals/54/documents/dd/issuances/dodd/300009p.pdf>, and any relevant successor guidance.

⁶ Although agencies are not required to apply the requirements of this memorandum to LLMs acquired pursuant to a free, open-source license, they should establish procedures to perform due diligence on the alignment of such models with E.O. 14319’s Unbiased AI Principles, as well as with the requirements of OMB Memorandum M-25-21, prior to use. Agencies should continue to responsibly leverage the benefits of open-source offerings, as directed by OMB Memorandum M-16-21.

⁷ E.O., sec. 4(a)(iv).

as a whole, rather than to test AI for a particular agency application;⁸ or

- d. AI used incidentally by a contractor during performance of a contract for administrative purposes (e.g., AI used at the option of a contractor when not directed or necessary to fulfill requirements).

Agency Actions

Agencies must comply with the following requirements, and take the additional implementation measures described in Appendix A:

a. Contractual Requirements for New Contracts

Agencies must ensure that any solicitation or order for procurement of an LLM they issue after the date of this memorandum includes contractual requirements addressing compliance with the Unbiased AI Principles, as explained in Appendix A. Agencies must also include such contractual requirements in any solicitation or order for AI models other than LLMs to which the requirements of this memorandum should apply, as determined by the agency.

b. Requirements for Existing Contracts

Agencies should, to the extent practicable, modify existing contracts for LLMs to include the requirements described in the preceding paragraph. Such a modification should occur, at the latest, prior to exercise of any option that extends the period of performance of the contract.

c. Updating Agency Procurement Policies

No later than March 11, 2026, agencies must update their policies and procedures to ensure that contracts for the procurement of LLMs include contractual requirements addressing compliance with the Unbiased AI Principles. Updated policies must include processes for agency users of LLMs to report outputs that violate the Unbiased AI Principles.

d. Sunset Provision

This memorandum shall cease to have any force or effect two years after the date of its issuance, unless the Director of OMB provides otherwise.

⁸ Examples include agency actions to develop, for general use, standards or testing methodologies for evaluating or red-teaming AI capabilities.

Appendix A: Implementation Guidance for Agencies

This guidance builds on existing directives and policies that address agencies' acquisition and use of AI, such as President Trump's Executive Order 13960, *Promoting the Use of Trustworthy Artificial Intelligence in the Federal Government*,⁹ and OMB Memoranda M-25-22, *Driving Efficient Acquisition of Artificial Intelligence*,¹⁰ and M-25-21, *Accelerating Federal Use of AI through Innovation, Governance, and Public Trust*.¹¹ By complying with those directives and policies as well as this memorandum, agencies will ensure that humans are accountable for agency actions and decisions involving LLMs.

1. CONTRACTUAL REQUIREMENTS

When procuring an LLM, agencies must obtain sufficient information from the vendor to determine whether that LLM complies with the Unbiased AI Principles. The amount and type of information available will vary depending on the vendor's role within the software supply chain and its relationship with the LLM developer itself, with more information generally being available from sources closer to the original LLM developer.

In Federal procurement, it is common for vendors to provide agencies with access to LLMs and other AI indirectly, through resellers or other intermediary deployers of an AI developer's product. Agencies also regularly contract with software integrators and platform operators who make LLMs available as part of their broader product offerings. When an agency transacts with such a third-party LLM provider, the availability of product information and potential for direct product interventions will depend on the willingness of the actual AI developer to collaborate through the third-party distributor. Agencies should consider these nuances of LLM procurement in determining how to apply the requirements of this memorandum to ensure that an LLM offered for procurement complies with the Unbiased AI Principles.¹²

Where practicable, agencies should avoid requirements that compel a vendor to disclose sensitive technical data, such as specific model weights. Documentation requests should seek enough information for an agency to assess a vendor's risk management actions at the model, system, and/or application level, as appropriate, to establish compliance with the Unbiased AI Principles.¹³

⁹ Executive Order 13960, Promoting the Use of Trustworthy Artificial Intelligence, (Dec. 3, 2020), <https://www.federalregister.gov/documents/2020/12/08/2020-27065/promoting-the-use-of-trustworthy-artificial-intelligence-in-the-federal-government>

¹⁰ OMB Memorandum M-25-22, Driving Efficient Acquisition of Artificial Intelligence in Government, (Apr. 3, 2025), <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf>

¹¹ OMB Memorandum M-25-21, *Accelerating Federal Use of AI Through Innovation, Governance, and Public Trust* (Apr. 3, 2025), <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-21-Accelerating-Federal-Use-of-AI-through-Innovation-Governance-and-Public-Trust.pdf>

¹² While the requirements of this memorandum focus on use of the disclosures to confirm sufficient compliance with the Unbiased AI Principles, such disclosures may also be used more broadly in the evaluation of offers, if an agency includes relevant evaluation factors as part of the solicitation.

¹³ In accordance with Sections 3(h) and 4(d)(G) of OMB Memorandum M-25-22, agencies should similarly consider requiring vendors to provide updated disclosures when integrating new AI enhancements, features, or components into the LLM under contract. Consistent with the requirements of Section 3(g) of OMB Memorandum M-25-22, best

In addition to obtaining information about standalone LLMs under consideration for procurement, agencies must also request information regarding LLM development and operation when those models are integrated into other software products or services to be procured.¹⁴

A. ***Minimum Threshold for LLM Transparency***

In solicitations for LLMs, agencies must request the following from a vendor:

i. Acceptable Use Policy.

This document is typically drafted by the original LLM developer to characterize and differentiate appropriate and inappropriate use of their product offering.

ii. Model, System, and/or Data Cards.

These materials from the LLM developer outline all essential information about the model, system, and/or data as it relates to the product offering. Information provided often includes summaries of the training process, identified risks and mitigations, and model evaluation scores on LLM benchmarks. It is uncommon for a vendor to produce all three types (model, system, data) of summary cards.

iii. End User Resources.

Such resources may include product tutorials, developer guides, or other best tools to help customers ensure proper use of the LLM and maximize utility.

iv. Mechanism For End User Feedback.

This requirement may be satisfied by a general inbox, specific product point of contact, or similar mechanism for providing feedback to the vendor on outputs that violate the Unbiased AI Principles.

B. ***Threshold for Enhanced LLM Transparency***

Depending on its planned use of a particular LLM, an agency may decide that it needs information beyond that identified above to assure itself that the model will align with the Unbiased AI Principles.¹⁵ Any request for such additional information should focus on actions that are directly relevant to the Principles, such as:

i. Pre-Training and Post-Training Activities.

- A. Actions undertaken that would impact the factuality and grounding of LLM outputs.
- B. System-level prompts that provide natural language instructions to the model specifying guidelines on responding to user-generated queries,

practices, including new sample AI transparency terms and conditions developed by GSA, will be made available in a shared repository, when available.

¹⁴ For example, the Unbiased AI Principle of “truth-seeking” will likely be more relevant than “ideological neutrality” for LLMs that operate as part of a tool to summarize cybersecurity incidents for a security operations center.

¹⁵ Agencies may determine that public-facing LLMs require additional vendor documentation to assess whether LLM outputs and associated controls or safeguards meet the goals of the Unbiased AI Principles.

particularly with regard to how a model responds when reliable information is incomplete, contradictory, or subject to individual interpretation.

- C. The type of outputs restricted via content moderation and safety filters.
- D. Use of red teaming as a means of continuously assessing the model to protect against incidents of bias in generated output.
- E. Any training or development that occurred outside the United States, to include the type of activity and the country in which it occurred.
- F. Any modifications or configurations undertaken to comply with any regulation from a government other than the U.S. Federal Government.

ii. Model Evaluations.

- A. Results of bias evaluations conducted by the vendor and the methodology for performing such tests (e.g., testing prompt pairs for politically-oriented topics).
- B. Benchmark scores for vendor evaluations to measure a model's bias, helpfulness, honesty, or accuracy when provided with ambiguous versus straightforward questions. In some instances, benchmarks to compare performance across multiple languages will be necessary.

iii. Enterprise-Level Controls.

- A. Governance tools, such as customizable system instructions that are additive to a base model's system prompts or content generation filters.
- B. LLM model evaluation tools that enable comparisons either between a model's outputs, or across various models.
- C. Product features that require a model to cite sources of its outputs or otherwise provide visibility into the provenance of the model's outputs.

iv. Third Party Modifications.

- A. Disclosure of additional controls to modify an LLMs output applied by the vendor if the vendor is not a direct developer of the LLM (e.g., classifiers, system prompts, fine-tuning, content moderation and filters).

C. Materiality Requirement

Agencies should explicitly identify relevant requirements identified in (a) and (b) as material to eligibility and payment under the contract, to support termination of the contract for default, as directed by the E.O. in cases where the vendor refuses to take corrective action in identified cases of noncompliance.

2. APPLICABILITY FACTORS

In determining whether to apply the requirements of this memorandum to LLMs developed by the agency, or to AI models other than LLMs, the agency shall consider:

A. *Models Developed by the Agency*

For use cases where an agency is developing an LLM or small language model, accompanying documentation that similarly demonstrates and assesses the integration of Unbiased AI Principles over the lifecycle of development is expected. While the precise nature of the types of information included may vary by use case, agencies should refer to the types of documentation described in Sections 1(a) and 1(b) of this Appendix to inform the development of their own documentation. At a minimum, agency documentation is expected to address:

- i. How the agency conducted the pre-training and training of the model;
- ii. How the agency proceeded to evaluate the model, and whether or how recurring or ongoing evaluations are performed;
- iii. What enterprise controls are built into the model, and how those controls may be accessed and configured; and
- iv. How the agency's development team considers feedback from end users of the LLM, to include relevant End User Resources as detailed in Section 1(a)(iii) of this Appendix.

B. *Models Other Than LLMs*

If the agency is procuring other types of generative AI capabilities, such as tools that assist with image, voice, or multimodal generation, agencies shall, where practicable, use this guidance to inform the documentation requirements imposed for the procurement.